

# MinhVo-FinalProject-MessageAnalysis

March 11, 2023

## 0.1 Big Data Platforms - Winter 2023

## 0.2 Final Project Education

### 0.2.1 Twitter - Unique Message Analysis

Minh Vo

```
[1]: import os
import time
import re
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from itertools import islice
# import sh
from pyspark.sql.functions import *
from pyspark.sql.types import *
from itertools import compress

pd.set_option('display.max_colwidth', None)
pd.reset_option('display.max_rows')
warnings.filterwarnings(action='ignore')
```

```
[2]: pip install -U nltk
```

Requirement already satisfied: nltk in /opt/conda/miniconda3/lib/python3.8/site-packages (3.8.1)

Requirement already satisfied: click in /opt/conda/miniconda3/lib/python3.8/site-packages (from nltk) (7.1.2)

Requirement already satisfied: tqdm in /opt/conda/miniconda3/lib/python3.8/site-packages (from nltk) (4.64.1)

Requirement already satisfied: regex>=2021.8.3 in /opt/conda/miniconda3/lib/python3.8/site-packages (from nltk) (2022.10.31)

Requirement already satisfied: joblib in /opt/conda/miniconda3/lib/python3.8/site-packages (from nltk) (1.2.0)

WARNING: Running pip as the 'root' user can result in broken permissions and conflicting behaviour with the system package manager. It is recommended to use a virtual environment instead: <https://pip.pypa.io/warnings/venv>

Note: you may need to restart the kernel to use updated packages.

```
[3]: import re
from pyspark.ml.feature import MinHashLSH
from pyspark.ml.feature import CountVectorizer, IDF, CountVectorizerModel, \
    Tokenizer, RegexTokenizer, StopWordsRemover
from pyspark import SparkContext
from pyspark.sql import SQLContext
from pyspark.sql import Row
import nltk
nltk.download('stopwords')
from nltk.corpus import stopwords
```

[nltk\_data] Downloading package stopwords to /root/nltk\_data...

[nltk\_data] Package stopwords is already up-to-date!

```
[4]: from google.cloud import storage
```

```
[5]: spark.conf.set("spark.sql.repl.eagerEval.enabled", True)
```

## 0.2.2 Data Loading

```
[6]: !hadoop fs -ls 'gs://msca-bdp-students-bucket/shared_data/mdvo/BDP-Final/
    ↳User-Groups/' | head -10
```

Found 341 items

```
-rwx----- 3 root root          0 2023-03-10 00:12 gs://msca-bdp-students-
bucket/shared_data/mdvo/BDP-Final/User-Groups/_SUCCESS
-rwx----- 3 root root 63332184 2023-03-10 00:06 gs://msca-bdp-students-
bucket/shared_data/mdvo/BDP-Final/User-
Groups/part-00000-b009d9a6-e554-48e0-a750-de696a114c43-c000.snappy.parquet
-rwx----- 3 root root 61951748 2023-03-10 00:06 gs://msca-bdp-students-
bucket/shared_data/mdvo/BDP-Final/User-
Groups/part-00001-b009d9a6-e554-48e0-a750-de696a114c43-c000.snappy.parquet
-rwx----- 3 root root 61932482 2023-03-10 00:06 gs://msca-bdp-students-
bucket/shared_data/mdvo/BDP-Final/User-
Groups/part-00002-b009d9a6-e554-48e0-a750-de696a114c43-c000.snappy.parquet
-rwx----- 3 root root 65444084 2023-03-10 00:06 gs://msca-bdp-students-
bucket/shared_data/mdvo/BDP-Final/User-
Groups/part-00003-b009d9a6-e554-48e0-a750-de696a114c43-c000.snappy.parquet
-rwx----- 3 root root 65033087 2023-03-10 00:06 gs://msca-bdp-students-
bucket/shared_data/mdvo/BDP-Final/User-
Groups/part-00004-b009d9a6-e554-48e0-a750-de696a114c43-c000.snappy.parquet
```

```
-rwx----- 3 root root 64988348 2023-03-10 00:07 gs://msca-bdp-students-
bucket/shared_data/mdvo/BDP-Final/User-
Groups/part-00005-b009d9a6-e554-48e0-a750-de696a114c43-c000.snappy.parquet
-rwx----- 3 root root 64801663 2023-03-10 00:06 gs://msca-bdp-students-
bucket/shared_data/mdvo/BDP-Final/User-
Groups/part-00006-b009d9a6-e554-48e0-a750-de696a114c43-c000.snappy.parquet
-rwx----- 3 root root 64615189 2023-03-10 00:06 gs://msca-bdp-students-
bucket/shared_data/mdvo/BDP-Final/User-
Groups/part-00007-b009d9a6-e554-48e0-a750-de696a114c43-c000.snappy.parquet
```

```
[7]: path = 'gs://msca-bdp-students-bucket/shared_data/mdvo/BDP-Final/User-Groups/'
```

```
[8]: %%time
```

```
twtgroup_df = spark.read.parquet(path)
twtgroup_df.limit(5)
```

23/03/10 22:47:05 WARN org.apache.spark.sql.catalyst.util.package: Truncated the string representation of a plan since it was too large. This behavior can be adjusted by setting 'spark.sql.debug.maxToStringFields'.

CPU times: user 5.08 ms, sys: 4.4 ms, total: 9.47 ms

Wall time: 7.3 s

```
[8]: +-----+-----+-----+-----+-----+-----+-----+-----+-----+
-----+-----+-----+-----+-----+-----+-----+-----+-----+
-----+-----+-----+-----+-----+-----+-----+-----+-----+
-----+-----+-----+-----+-----+-----+-----+-----+-----+
-----+-----+-----+-----+-----+-----+-----+-----+-----+
-----+
|coordinates|favorite_count|filter_level|in_reply_to_screen_name|retweeted|retwe
et_count|retweeted_from|    retweeted_status|
text|country|country_code|full_name|place_type|bounding_box| timestamp_ms|
id_str|          name|    screen_name|    location|
description|followers_count|statuses_count|
created_at|verified|lang|is_original_twt|Twitterer_Group|
+-----+-----+-----+-----+-----+-----+-----+-----+
-----+-----+-----+-----+-----+-----+-----+-----+-----+
-----+-----+-----+-----+-----+-----+-----+-----+-----+
-----+-----+-----+-----+-----+-----+-----+-----+-----+
-----+-----+-----+-----+-----+-----+-----+-----+-----+
-----+
|      null|          0|      low|          null|      RT|
0| David_Moscrop|{null, Fri May 13...|RT @David_Moscrop...|    null|      null|
null|      null|      null|1652469985395|1427351911410257934|
Mike|    Mike24084394|      null|Long time NDP mem...|      122|
```

```

2244|Mon Aug 16 19:29:...| false| en| 0| Others|
| null| 0| low| null| RT|
0| NYANGBERRY|{null, Fri May 13...|RT @NYANGBERRY: A...| null| null|
null| null| null|1652469987924| 959799880322637825| æ-Dani
(. )| JaemNoJeno| she/her|Dreams Come True
...| 2521| 29441|Sat Feb 03 14:44:...| false| en|
0| Others|
| null| 0| low| null| RT|
0| realchrisrufo|{null, Fri May 13...|RT @realchrisrufo...| null| null|
null| null| null|1652469990292|1104477873270898688|
Jessthemes|JesstheMesstake| USA|If the road to he...|
223| 15102|Sat Mar 09 20:23:...| false| en| 0|
Others|
| null| 0| low| MattOswaltVA| |
0| null| null|@MattOswaltVA I l...| null| null|
null| null| null|1652469990754|1491083841838239745| Todd
Weaver| DreamWeaver_tw|Los Angeles, CA|Artist/Photograph...| 7|
68|Tue Feb 08 16:18:...| false| en| 1| Others|
| null| 0| low| CoachFelecia| |
0| null| null|@CoachFelecia doe...| null| null|
null| null| null|1652469991702| 832788979393060867|Mrs_Pinky
Nic...| MRSpinkston85| Georgia, USA|Christian| Wife &...| 14863|
217346|Sat Feb 18 03:08:...| false| en| 1| Others|
+-----+-----+-----+-----+-----+-----+-----+
-----+-----+-----+-----+-----+-----+-----+
-----+-----+-----+-----+-----+-----+-----+
-----+-----+-----+-----+-----+-----+-----+
-----+-----+-----+-----+-----+-----+-----+
-----+

```

```
[9]: twtgroup_df.count()
```

```
[9]: 24721729
```

```
[10]: twtgroup_df.groupby('Twitterer_Group').count().orderBy("count",ascending=False).
      ↳toPandas()
```

```
[10]:
```

	Twitterer_Group	count
0	Others	24343776
1	Social Media Influencers	226766
2	News Outlets	127342
3	Universities	9641
4	Schools	6209

5	Government Entities	5067
6	Non-Profit Organizations	2928

### 0.2.3 How unique are the messages?

- Are they mostly unique? Or usually people are just copy-pasting the same text?
- You can use something like Jaccard similarity / Cosine Similarity / Simhash / Minhash to measure uniqueness / similarity
- Visualize message duplication for each group of Twitterers (government entities / non-profit organizations / news outlets / social media influencers / other) *Please note: this is not a topic modeling (LDA / LSA) – but text similarity analysis*

```
[11]: # work on Original tweets only
original_twts_group = twtgroup_df.filter(col('is_original_twt')==1)
```

```
[12]: original_twts_group.count()
```

```
[12]: 8174910
```

```
[13]: original_twts_group.groupby('Twitterer_Group').count().
      ↳orderBy("count",ascending=False).toPandas()
```

```
[13]:
```

	Twitterer_Group	count
0	Others	7933421
1	Social Media Influencers	129547
2	News Outlets	97068
3	Universities	6315
4	Schools	3852
5	Government Entities	2764
6	Non-Profit Organizations	1943

```
[14]: nonprofit_twt = original_twts_group.filter(col('Twitterer_Group')== 'Non-Profit_
      ↳Organizations')
gov_twt = original_twts_group.filter(col('Twitterer_Group')== 'Government_
      ↳Entities')
univ_twt = original_twts_group.filter(col('Twitterer_Group')== 'Universities')
school_twt = original_twts_group.filter(col('Twitterer_Group')== 'Schools')
news_twt = original_twts_group.filter(col('Twitterer_Group')== 'News Outlets')
influencer_twt = original_twts_group.filter(col('Twitterer_Group')== 'Social_
      ↳Media Influencers')
other_twt = original_twts_group.filter(col('Twitterer_Group')== 'Others')
```

```
[15]: group_sample = tweets_sample.select('id_str','text','Twitterer_Group')
group_sample.limit(5).toPandas()
```

```
-----
NameError                                Traceback (most recent call last)
Cell In[15], line 1
----> 1 group_sample = tweets_sample.select('id_str','text','Twitterer_Group')
      2 group_sample.limit(5).toPandas()

NameError: name 'tweets_sample' is not defined
```

```
[ ]: # nonprofit_text = group_sample.filter(col('Twitterer_Group')== 'Non-Profit_
↳ Organizations')
# gov_text = group_sample.filter(col('Twitterer_Group')== 'Government Entities')
# univ_text = group_sample.filter(col('Twitterer_Group')== 'Universities')
# school_text = group_sample.filter(col('Twitterer_Group')== 'Schools')
# news_text = group_sample.filter(col('Twitterer_Group')== 'News Outlets')
# influencer_text = group_sample.filter(col('Twitterer_Group')== 'Social Media_
↳ Influencers')
# other_text = group_sample.filter(col('Twitterer_Group')== 'Others')
```

### Create a function to run the jaccard similarity

```
[23]: def generate_jaccard_sim(df, jaccard_distance):
    text = df.rdd.map(lambda x : x['text']).filter(lambda x: x is not None)

    # remove stopwords and tokenize text
    Stop_Words = stopwords.words("english")
    tokens = text\
        .map( lambda document: document.strip().lower())\
        .map( lambda document: re.split(r"\s+", document))\
        .map( lambda word: [x for x in word if x.isalnum()])\
        .map( lambda word: [x for x in word if len(x) > 3] )\
        .map( lambda word: [x for x in word if x not in Stop_Words])\
        .zipWithIndex()

    row = Row('text')
    df_text = text.map(row).zipWithIndex().toDF(['text','id'])

    #Drop records with no tokens
    df_tokens = spark.createDataFrame(tokens, ["list_of_words", "id"])
    df_tokens = df_tokens.where(col('list_of_words').getItem(0).isNotNull())

    # vectorize tokens using CountVectorizer
    vectorize = CountVectorizer(inputCol="list_of_words", outputCol="features",
↳ minDF=1.0)
```

```

df_vectorize = vectorize.fit(df_tokens).transform(df_tokens)

# hash vectorized tokens using MinHashLSH
mh = MinHashLSH(inputCol="features", outputCol="hashes", numHashTables=5)
model = mh.fit(df_vectorize)
df_hashed = mh.fit(df_vectorize).transform(df_vectorize).cache()

df_hashed_text = df_text.join(df_hashed, "id", how='left').cache()

# apply approximate Jaccard similarity on hashed tokens
dups_text = model.approxSimilarityJoin(df_hashed_text, df_hashed_text,
→jaccard_distance).filter("datasetA.id < datasetB.id").select(
    col("distCol"),
    col("datasetA.id").alias("id_A"),
    col("datasetB.id").alias("id_B"),
    col('datasetA.text').alias('text_A'),
    col('datasetB.text').alias('text_B'))
dups_text.cache()

# calculate number of unique and total records
dups = dups_text.select('id_A').distinct().count()
records = df_hashed_text.count()
uniques = records - dups

return {'near_dups': dups, 'unique': uniques, 'records': records}

```

**Pick a random sample to run the text similarity**

```

[53]: # Run on original tweets only
# Choose a random sample of 1% from the original tweets data for analysis to
→avoid the crashing issue (due to the huge amount of data).
tweets_sample = original_twts_group.sample(fraction=0.01)
tweets_sample.count()

```

[53]: 81809

```

[54]: tweets_sample.groupby('Twitterer_Group').count().
→orderBy("count",ascending=False).toPandas()

```

```

[54]:
      Twitterer_Group  count
0              Others  79393
1  Social Media Influencers   1329
2              News Outlets    930

```

3	Universities	77
4	Schools	38
5	Government Entities	25
6	Non-Profit Organizations	17

```
[56]: text_sample = tweets_sample.select('id_str','text','Twitterer_Group')
text_sample_similarity = generate_jaccard_sim(text_sample, 0.3)
text_sample_similarity
```

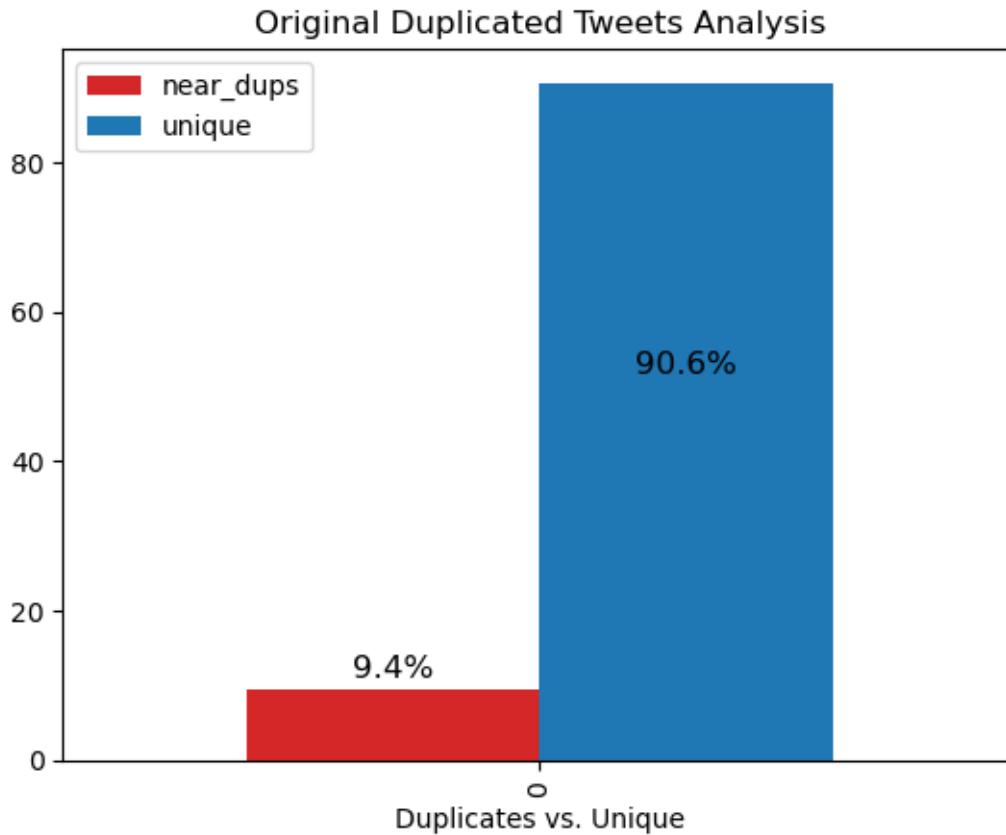
```
[56]: {'near_dups': 7689, 'unique': 74120, 'records': 81809}
```

```
[59]: dups = text_sample_similarity['near_dups']/text_sample_similarity['records']*100
uniques = text_sample_similarity['unique']/text_sample_similarity['records']*100

text_sample_similarity_df = pd.DataFrame.from_dict({'near_dups': [dups],
→ 'unique': [uniques]})

ax = text_sample_similarity_df.plot(kind = 'bar',y=['near_dups', 'unique'],
→ fontsize=10, color=['C3', 'C0'], align='center', width=0.8,
→ xlabel="Duplicates vs. Unique")
ax.set_title('Original Duplicated Tweets Analysis', fontsize=12)
for p in ax.patches:
    ax.annotate(f"{format(p.get_height(), '.1f')}%",
                (p.get_x() + p.get_width() / 2., p.get_height()/2),
                ha = 'center', va = 'center',
                xytext = (0,21),
                textcoords = 'offset points',
                fontsize = 12)
```





Run the text similarity on each user group

**Non-profit organizations**

```
[ ]: nonprofit_twt = original_twts_group.filter(col('Twitterer_Group')== 'Non-Profit_
↳ Organizations').select('id_str', 'text', 'Twitterer_Group')
nonprofit_similarity = generate_jaccard_sim(nonprofit_twt, 0.3)
nonprofit_similarity
```

```
[ ]: {'near_dups': 303, 'unique': 1640, 'records': 1943}
```

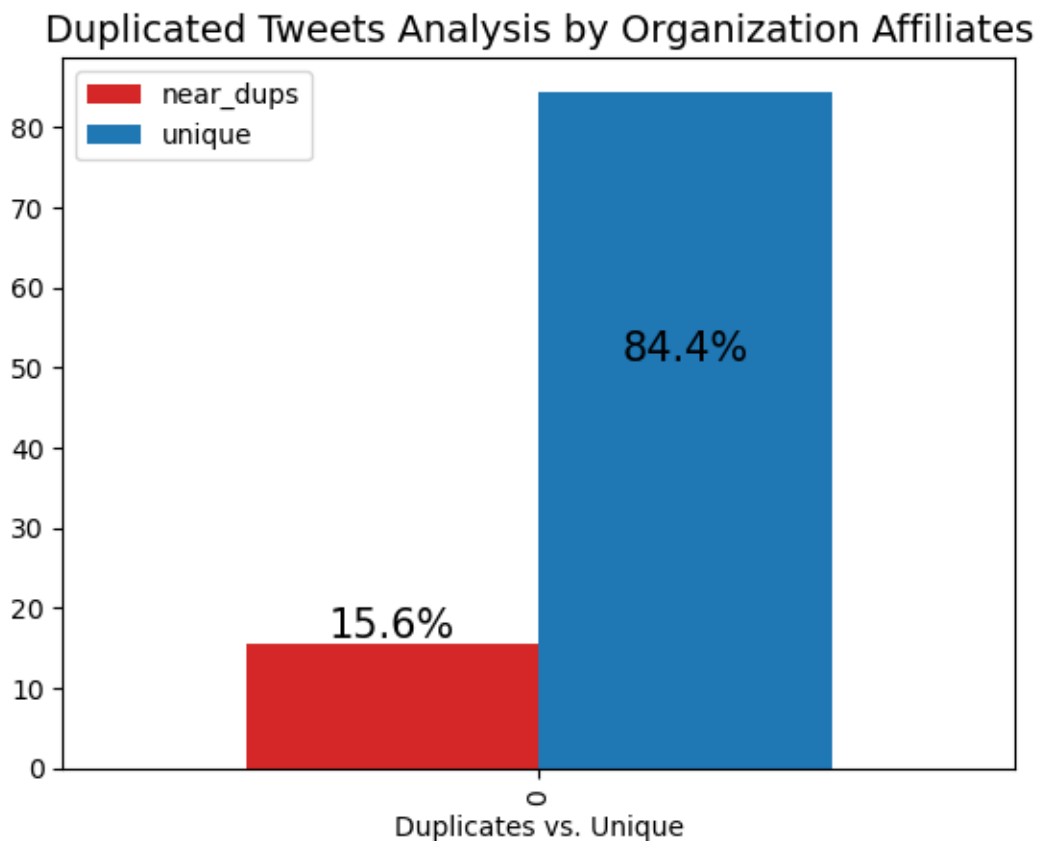
```
[63]: dups = nonprofit_similarity['near_dups']/nonprofit_similarity['records']*100
        uniques = nonprofit_similarity['unique']/nonprofit_similarity['records']*100

        nonprofit_similarity_df = pd.DataFrame.from_dict({'near_dups': [dups], 'unique':
↳ [uniques]})
```

```

ax = nonprofit_similarity_df.plot(kind = 'bar',y=['near_dups', 'unique'],
    ↳ fontsize=10, color=['C3', 'C0'], align='center', width=0.8,
    ↳ xlabel="Duplicates vs. Unique")
ax.set_title('Duplicated Tweets Analysis by Organization Affiliates',
    ↳ fontsize=14)
for p in ax.patches:
    ax.annotate(f"{format(p.get_height(), '.1f')}%",
                (p.get_x() + p.get_width() / 2., p.get_height()/2),
                ha = 'center', va = 'center',
                xytext = (0,30),
                textcoords = 'offset points',
                fontsize = 15)

```



### Government Entities

```

[64]: gov_twt = original_twts_group.filter(col('Twitterer_Group')=='Government_
    ↳ Entities').select('id_str', 'text', 'Twitterer_Group')
gov_similarity = generate_jaccard_sim(gov_twt, 0.3)
gov_similarity

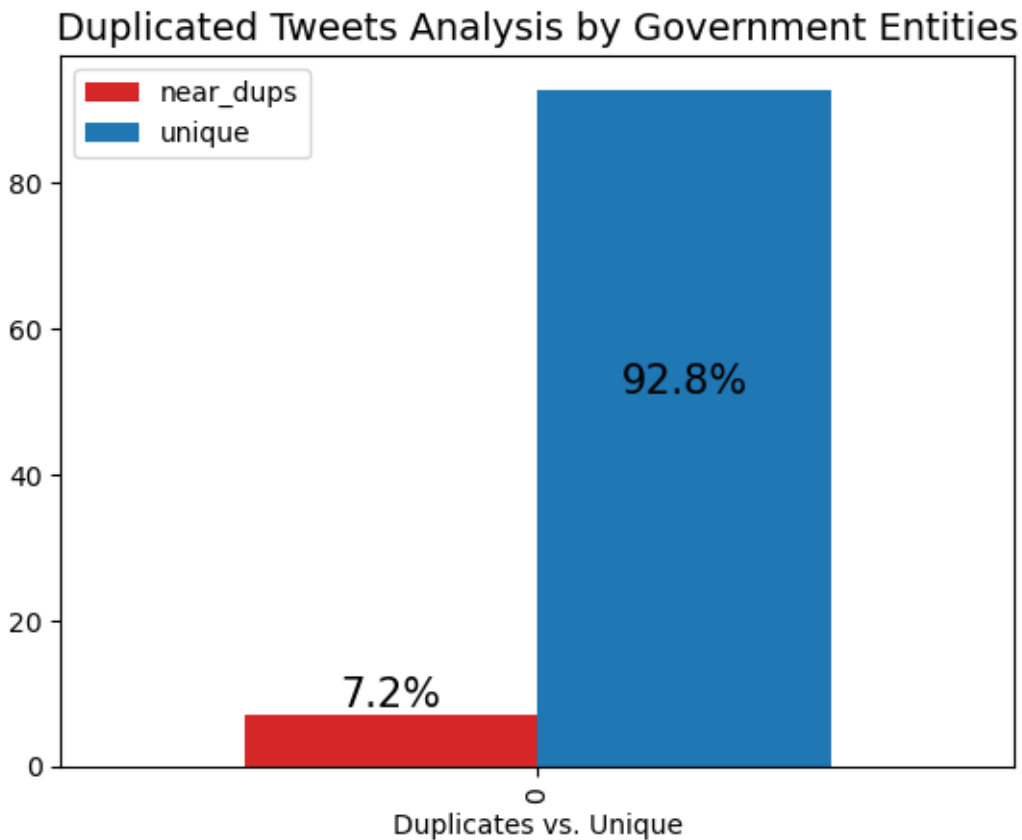
```

```
[64]: {'near_dups': 199, 'unique': 2565, 'records': 2764}
```

```
[65]: dups = gov_similarity['near_dups']/gov_similarity['records']*100
      uniques = gov_similarity['unique']/gov_similarity['records']*100

      gov_similarity_df = pd.DataFrame.from_dict({'near_dups': [dups], 'unique': [uniques]})

      ax = gov_similarity_df.plot(kind = 'bar', y=['near_dups', 'unique'],
      ↪ fontsize=10, color=['C3', 'C0'], align='center', width=0.8,
      ↪ xlabel="Duplicates vs. Unique")
      ax.set_title('Duplicated Tweets Analysis by Government Entities', fontsize=14)
      for p in ax.patches:
          ax.annotate(f"{format(p.get_height(), '.1f')}%",
                      (p.get_x() + p.get_width() / 2., p.get_height()/2),
                      ha = 'center', va = 'center',
                      xytext = (0,17),
                      textcoords = 'offset points',
                      fontsize = 15)
```

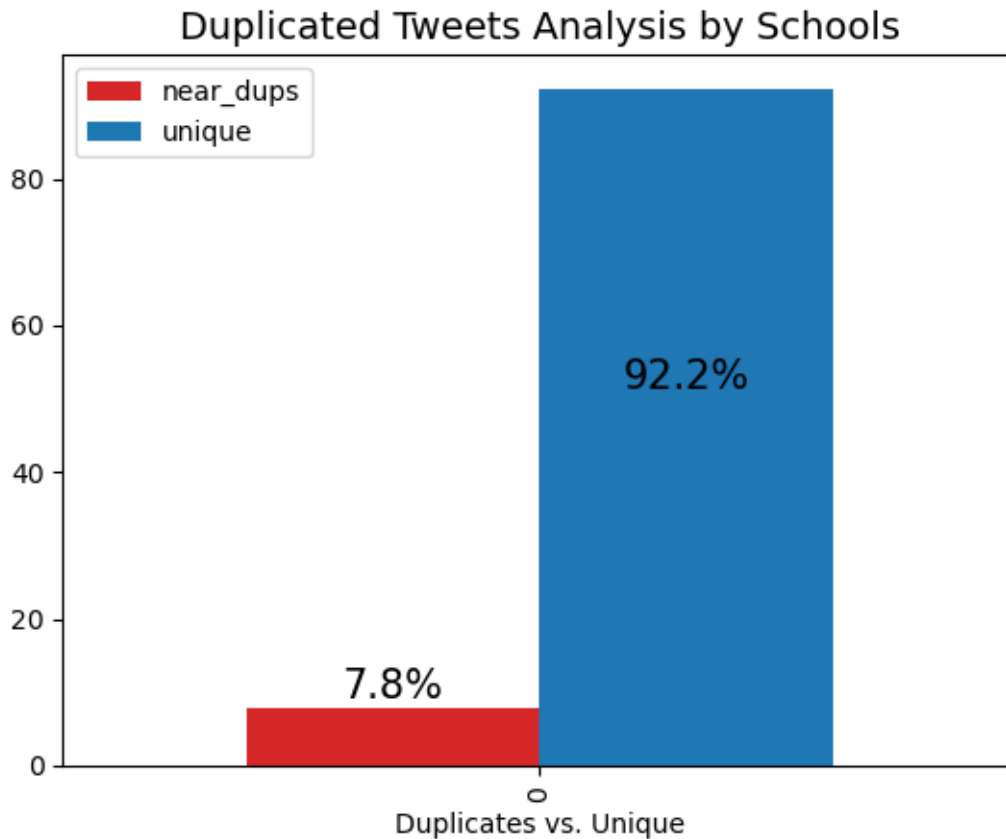


## Schools

```
[32]: school_twt = original_twts_group.filter(col('Twitterer_Group')== 'Schools').  
      ↪select('id_str','text','Twitterer_Group')  
      school_similarity = generate_jaccard_sim(school_twt, 0.3)  
      school_similarity
```

```
[32]: {'near_dups': 299, 'unique': 3553, 'records': 3852}
```

```
[66]: dups = school_similarity['near_dups']/school_similarity['records']*100  
      uniques = school_similarity['unique']/school_similarity['records']*100  
  
      school_similarity_df = pd.DataFrame.from_dict({'near_dups': [dups], 'unique':  
      ↪[uniques]})  
  
      ax = school_similarity_df.plot(kind = 'bar',y=['near_dups', 'unique'],  
      ↪fontsize=10, color=['C3', 'C0'], align='center', width=0.8,  
      ↪xlabel="Duplicates vs. Unique")  
      ax.set_title('Duplicated Tweets Analysis by Schools', fontsize=14)  
      for p in ax.patches:  
          ax.annotate(f"{format(p.get_height(), '.1f')}%",  
                      (p.get_x() + p.get_width() / 2., p.get_height()/2),  
                      ha = 'center', va = 'center',  
                      xytext = (0,19),  
                      textcoords = 'offset points',  
                      fontsize = 15)
```



#### Universities

```
[36]: univ_twt = original_twts_group.filter(col('Twitterer_Group')== 'Universities').
      ↳select('id_str', 'text', 'Twitterer_Group')
univ_similarity = generate_jaccard_sim(univ_twt, 0.3)
univ_similarity
```

```
[36]: {'near_dups': 423, 'unique': 5892, 'records': 6315}
```

```
[38]: dups = univ_similarity['near_dups']/univ_similarity['records']*100
      uniques = univ_similarity['unique']/univ_similarity['records']*100

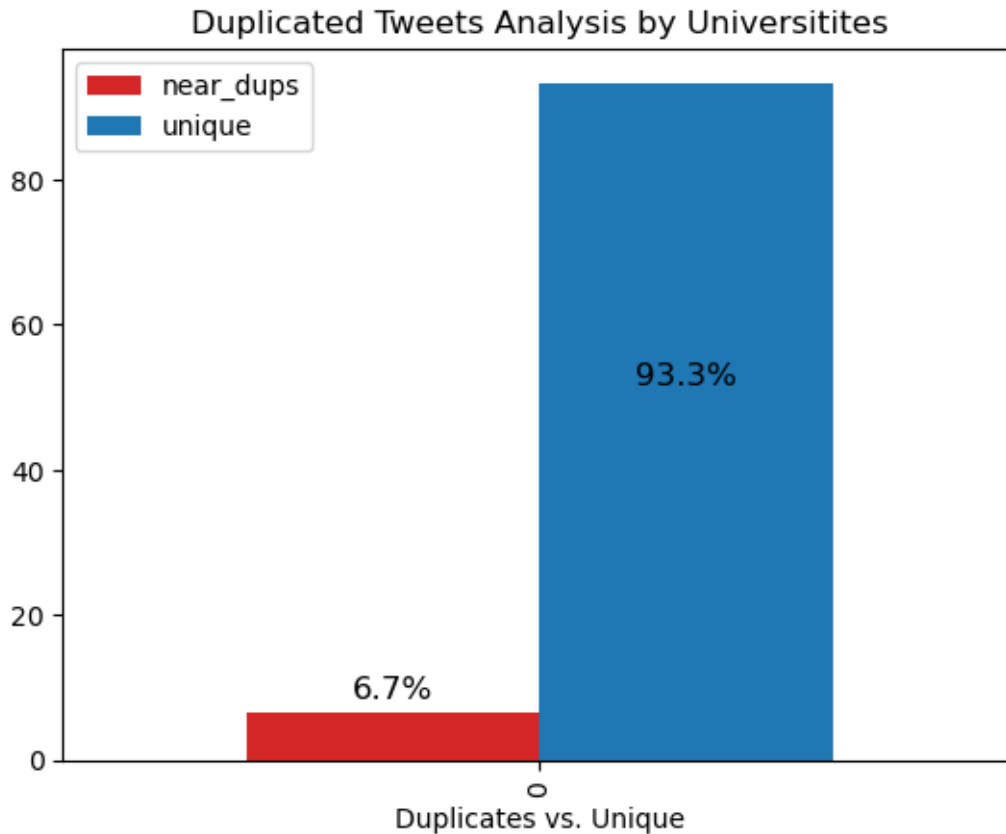
      univ_similarity_df = pd.DataFrame.from_dict({'near_dups': [dups], 'unique': [
      ↳uniques]})

      ax = univ_similarity_df.plot(kind = 'bar', y=['near_dups', 'unique'],
      ↳fontsize=10, color=['C3', 'C0'], align='center', width=0.8,
      ↳xlabel="Duplicates vs. Unique")
```

```

ax.set_title('Duplicated Tweets Analysis by Universitites', fontsize=12)
for p in ax.patches:
    ax.annotate(f"{format(p.get_height(), '.1f')}%",
                (p.get_x() + p.get_width() / 2., p.get_height()/2),
                ha = 'center', va = 'center',
                xytext = (0,17),
                textcoords = 'offset points',
                fontsize = 12)

```



### News Outlet

```

[39]: news_twt = original_twt_group.filter(col('Twitterer_Group')=='News Outlets').
      ↪select('id_str','text','Twitterer_Group')
news_twt_sample = news_twt.sample(fraction=0.1)
news_similarity = generate_jaccard_sim(news_twt_sample, 0.3)
news_similarity

```

```


[39]: {'near_dups': 361, 'unique': 9227, 'records': 9588}

```

```
[69]: dups = news_similarity['near_dups']/news_similarity['records']*100
        uniques = news_similarity['unique']/news_similarity['records']*100

        news_similarity_df = pd.DataFrame.from_dict({'near_dups': [dups], 'unique': [uniques]})

        ax = news_similarity_df.plot(kind = 'bar',y=['near_dups', 'unique'],
        ↪fontsize=10, color=['C3', 'C0'], align='center', width=0.8,
        ↪xlabel="Duplicates vs. Unique")
        ax.set_title('Duplicated Tweets Analysis by News Outlet', fontsize=14)
        for p in ax.patches:
            ax.annotate(f"{format(p.get_height(), '.1f')}%",
                        (p.get_x() + p.get_width() / 2., p.get_height()/2),
                        ha = 'center', va = 'center',
                        xytext = (0,12),
                        textcoords = 'offset points',
                        fontsize = 15)
```



Social Media Influencer

```
[41]: influencer_twt = original_twts_group.filter((col('Twitterer_Group')== 'Social_Media Influencers') & (col('verified')==True)).
      ↪select('id_str','text','Twitterer_Group')
influencer_twt_sample = influencer_twt.sample(fraction=0.1)
influencer_similarity = generate_jaccard_sim(influencer_twt_sample, 0.3)
influencer_similarity
```

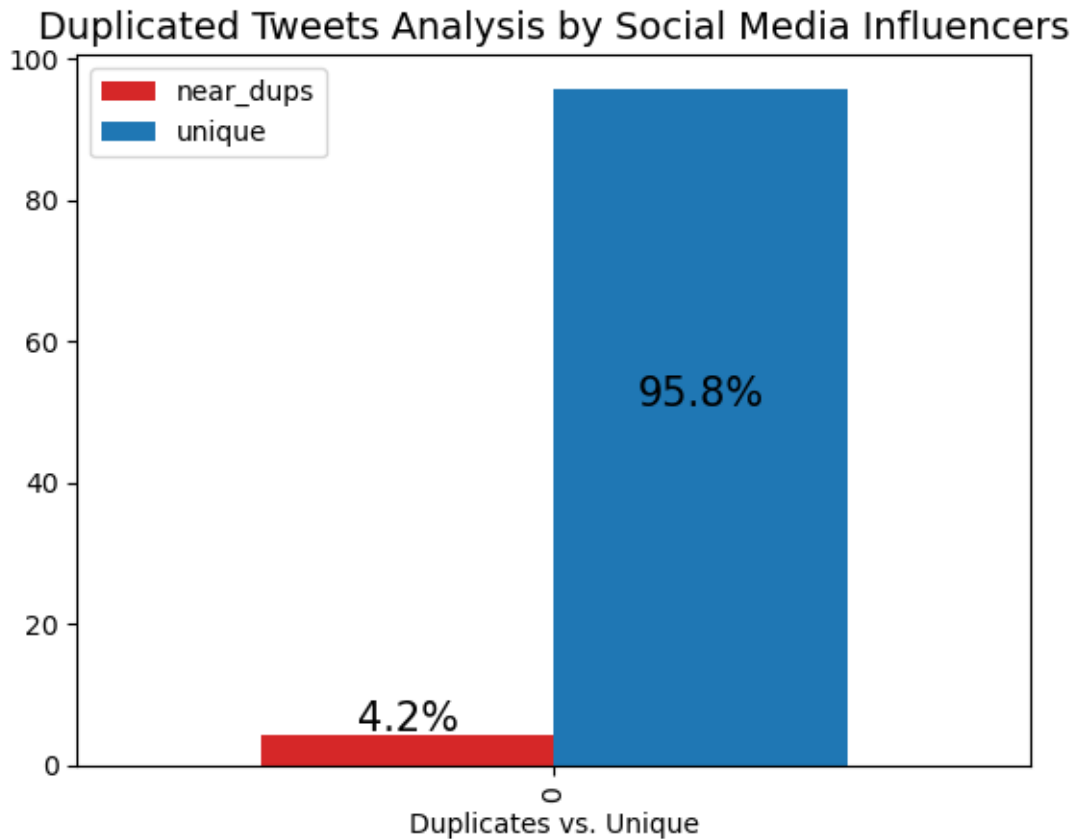
```
[41]: {'near_dups': 547, 'unique': 12410, 'records': 12957}
```

```
[70]: dups = influencer_similarity['near_dups']/influencer_similarity['records']*100
      uniques = influencer_similarity['unique']/influencer_similarity['records']*100

influencer_similarity_df = pd.DataFrame.from_dict({'near_dups': [dups],
      ↪'unique': [uniques]})

ax = influencer_similarity_df.plot(kind = 'bar',y=['near_dups', 'unique'],
      ↪fontsize=10, color=['C3', 'C0'], align='center', width=0.8,
      ↪xlabel="Duplicates vs. Unique")
ax.set_title('Duplicated Tweets Analysis by Social Media Influencers',
      ↪fontsize=14)
for p in ax.patches:
    ax.annotate(f"{format(p.get_height(), '.1f')}%",
                (p.get_x() + p.get_width() / 2., p.get_height()/2),
                ha = 'center', va = 'center',
                xytext = (0,12),
                textcoords = 'offset points',
                fontsize = 15)
```





#### Others/Someone Else

```
[43]: other_twt = original_twts_group.filter(col('Twitterer_Group')== 'Others').
      ↪select('id_str','text','Twitterer_Group')
other_twt_sample = other_twt.sample(fraction=0.005)
other_similarity = generate_jaccard_sim(other_twt_sample, 0.3)
other_similarity
```

```
[43]: {'near_dups': 3190, 'unique': 36300, 'records': 39490}
```

```
[ ]: other_similarity = generate_jaccard_sim(other_twt_sample, 0.3)
other_similarity
```

```
[71]: dups = other_similarity['near_dups']/other_similarity['records']*100
      uniques = other_similarity['unique']/other_similarity['records']*100

      other_similarity_df = pd.DataFrame.from_dict({'near_dups': [dups], 'unique': [uniques]})
      ↪[uniques])
```

```

ax = other_similarity_df.plot(kind = 'bar',y=['near_dups', 'unique'],
    ↳ fontsize=10, color=['C3', 'C0'], align='center', width=0.8,
    ↳ xlabel="Duplicates vs. Unique")
ax.set_title('Duplicated Tweets Analysis by Other Users', fontsize=14)
for p in ax.patches:
    ax.annotate(f"{format(p.get_height(), '.1f')}}%",
                (p.get_x() + p.get_width() / 2., p.get_height()/2),
                ha = 'center', va = 'center',
                xytext = (0,17),
                textcoords = 'offset points',
                fontsize = 15)

```

